


Lecture 5: Transactions

Olivier Liechti
AMT

heig-vd


Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud



Transactions

Transactions

Imagine that you have the following code in a business service:

 `accountA.debit(100);`
`accountB.credit(100);`

What happens if the application crashes here?
Is my data corrupted?
Has money vanished in cyberspace?

transaction.start();

accountA.debit(100);
accountB.credit(100);

transaction.commit();

Transactions give us an “whole or nothing” semantic
(we often speak about a unit of work)

```
transaction.start();  
accountA.debit(100);  
try {  
    accountB.credit(100);  
} catch (AccountFullException e) {  
    transaction.rollback();  
}  
transaction.commit();
```

We can also deal with application-level errors and leave the data in a consistent state.

ACID

ACID

Atomicity: “all or nothing”

ACID

Consistency: “business data integrity”


ACID

Isolation: “deal with concurrent transactions”

There are different isolation levels!

Isolation levels

Increasing isolation
between
transactions



Isolation level	Potential issues
Read Uncommitted (no locks)	Dirty Reads (no isolation)
Read Committed (write locks)	Non-repeatable Reads
Repeatable Reads (read & write locks)	Phantom reads
Serializable (range locks)	

Increasing performance in the
cas of concurrent access

Isolation levels

- "A **dirty read** occurs when a transaction is allowed to read data from a row that has been modified by another running transaction and not yet committed."
- "A **non-repeatable read** occurs, when during the course of a transaction, a row is retrieved twice and the values within the row differ between reads."
- "A **phantom read** occurs when, in the course of a transaction, two identical (SELECT) queries are executed, and the **collection** of rows returned by the second query is different from the first."

See for **example scenarios**, see:

https://docs.oracle.com/javase/tutorial/jdbc/basics/transactions.html#transactions_data_integrity


[https://en.wikipedia.org/wiki/Isolation_\(database_systems\)#Read_phenomena](https://en.wikipedia.org/wiki/Isolation_(database_systems)#Read_phenomena)

ACID


Durability: “once it’s done, it’s done”

Transactions & EJBs

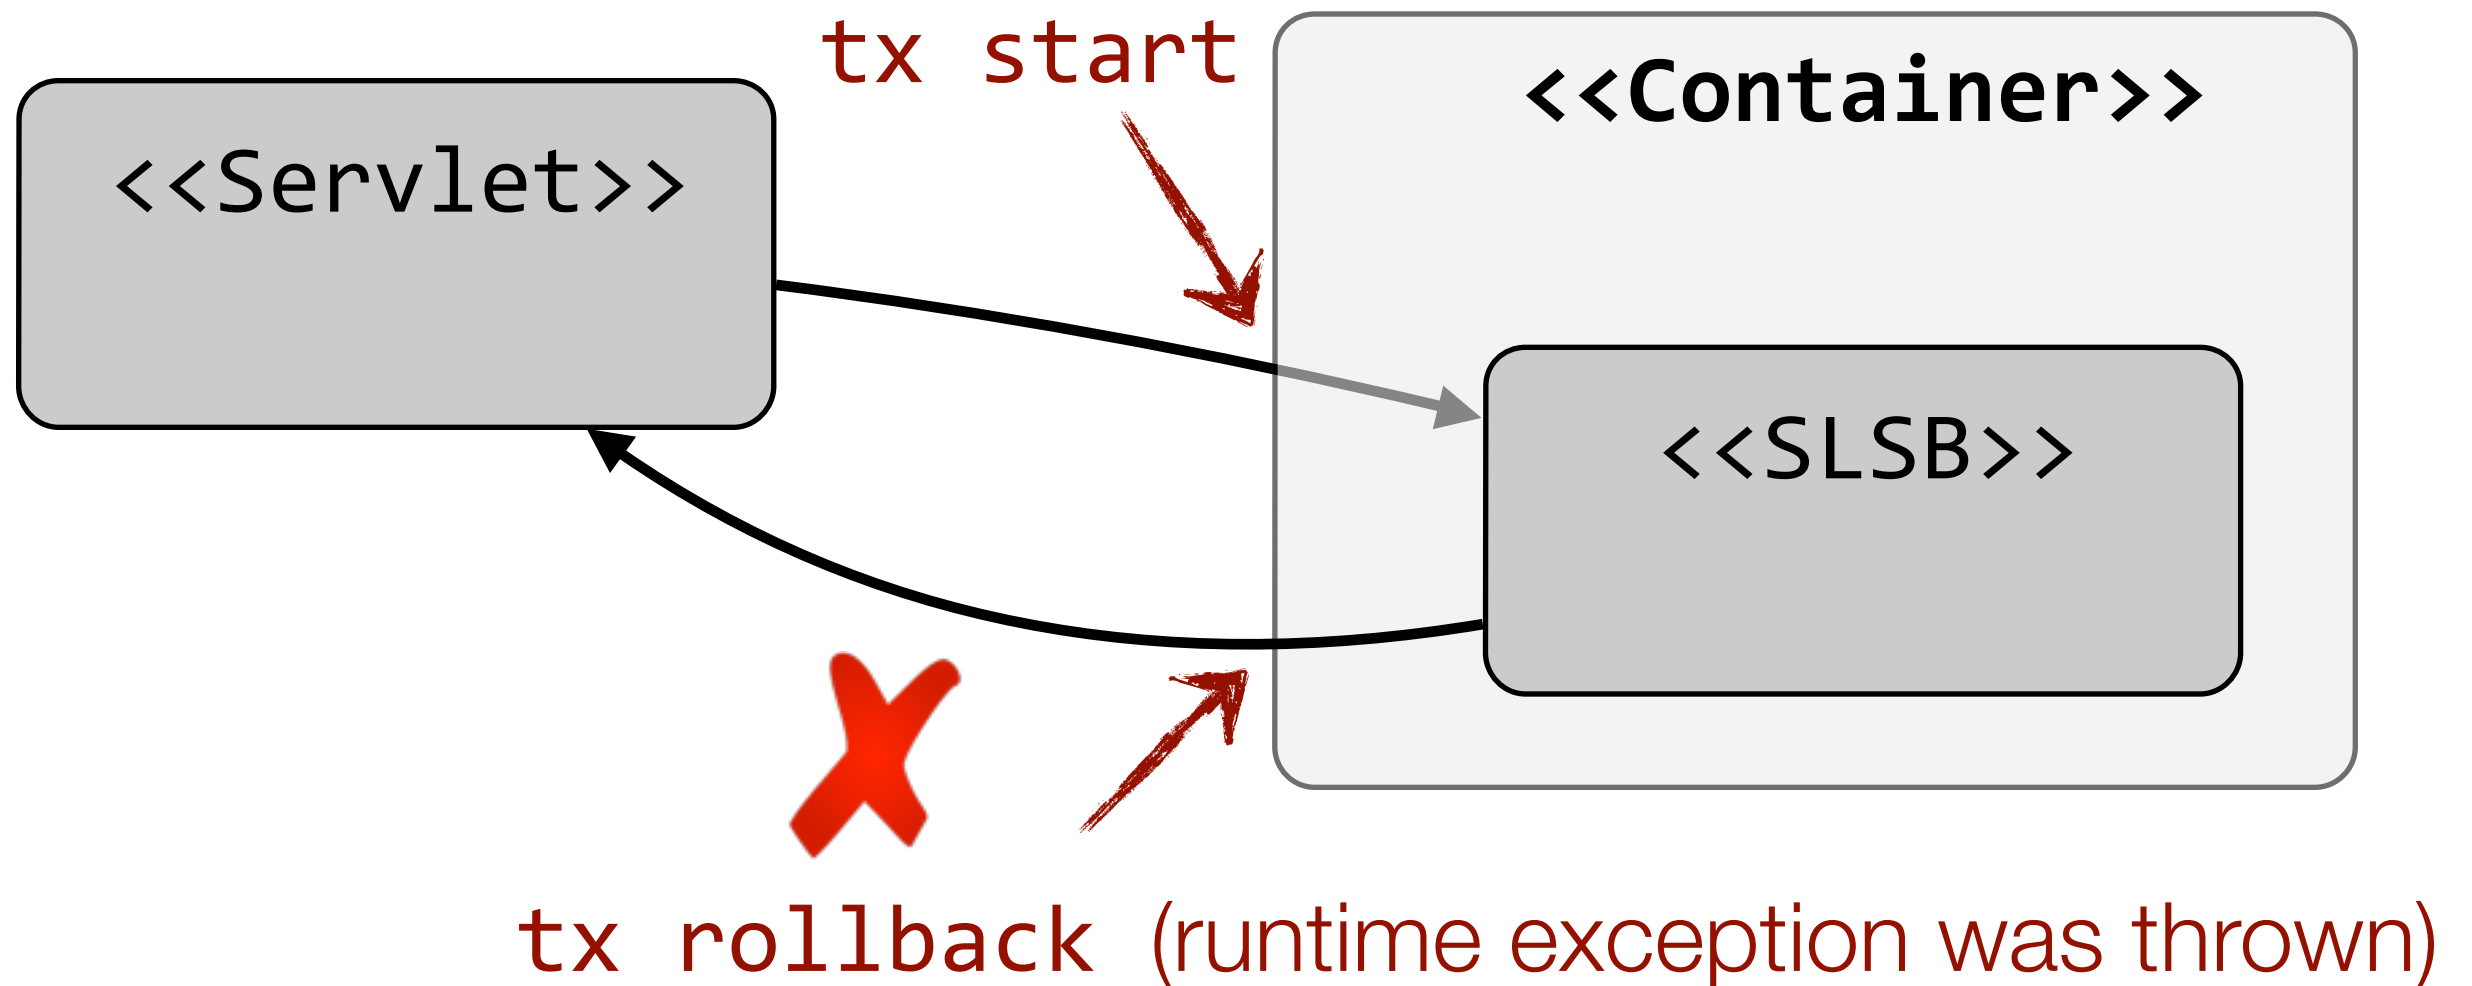
- By default, the EJB container handles calls to `commit` and `rollback`.
- Methods defined on EJBs provide demarcation points.
- This is the **default behavior**.



Transactions & EJBs



Transactions & EJBs



What happens when a **client** calls a method on a session bean,

which calls a method on a session bean,
which calls a method on a session bean,
which calls a method on a session bean,
which calls a method on a session bean,
which calls a method on a session bean,
which calls a method on a session bean,

which **throws an exception?**

Transaction Scope

What happens when a **client** calls a method on a session bean,

Opinion1

Everything should be rolled back!

which calls a method on a session bean,
which calls a method on a session bean,

which calls a method on a session bean,

which calls a method on a session bean,

which calls a method on a session bean,

Opinion2

No! Only changes incurred by the last method should be rolled back!

which **throws an exception?**

Transaction Scope

What happens when a **client** calls a method on a session?

Opinion1

Everything should be rolled back!


It is **up to the application** to specify intended behavior. The developer must specify transaction scope, typically with **annotations**.

Opinion2

No! Only changes incurred by the last method should be rolled back!

which **throws an exception**?

Transaction Scope



<http://java.sun.com/javaee/5/docs/tutorial/doc/bncij.html>

Transaction Scope

heig-vd


Haute Ecole d'Ingénierie et de Gestion
du Canton de Vaud

```
@TransactionAttribute(NOT_SUPPORTED)
@Stateless
public class TransactionBean implements
Transaction {
...
    @TransactionAttribute(REQUIRES_NEW)
    public void firstMethod() {...}

    @TransactionAttribute(REQUIRED)
    public void secondMethod() {...}

    public void thirdMethod() {...}

    public void fourthMethod() {...}
}
```



Transaction Attribute	Client's Transaction	Business Method's Transaction
Required	None	T2
	T1	T1
RequiresNew	None	T2
	T1	T2
Mandatory	None	error
	T1	T1
NotSupported	None	None
	T1	None
Supports	None	None
	T1	T1
Never	None	None
	T1	Error

Transactions & Exceptions

- There are **two ways to roll back a container-managed transaction**
- Firstly, if a **system exception** is thrown, the container will automatically roll back the transaction.
- Secondly, by invoking the **setRollbackOnly** method of the EJBContext interface, the bean method instructs the container to roll back the transaction.
- If the bean throws an **application exception**, the rollback is not automatic but can be initiated by a call to **setRollbackOnly**.
- Note: you can also annotate your Exception class with **@ApplicationException(rollback=true)**

Transaction scope & JPA

What happens if there is a strike and a **NullPointerException** is thrown in the constructor?

```
@Stateless
public class CarService {

    @PersistenceContext
    EntityManager em;

    @EJB
    PartsService partsService;

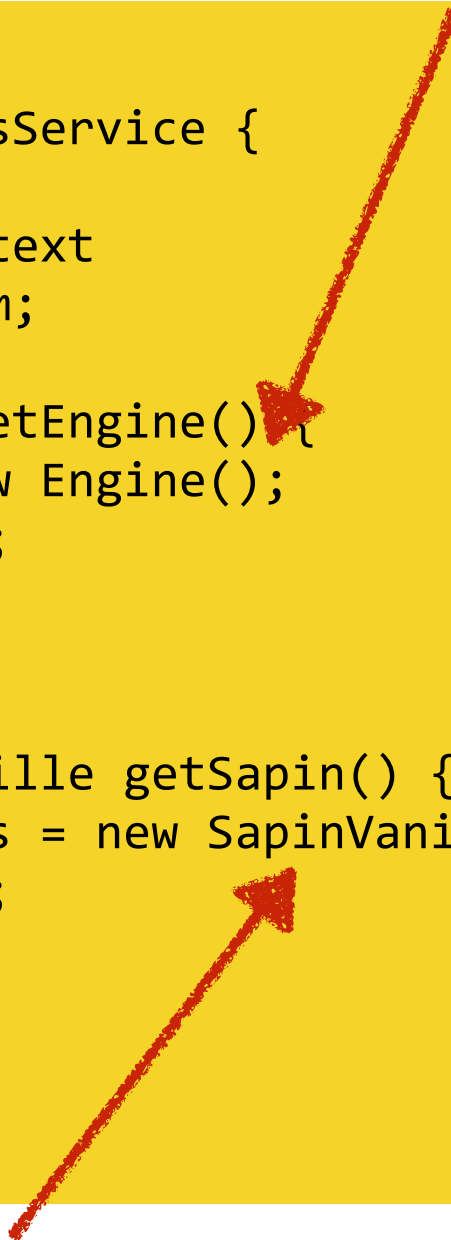
    public Car buildCar() {
        Engine e = getEngine();
        SapinVanille s = getSapin();
        Car c = new Car(e, s);
        em.persist(c);
    }
}
```

```
@Stateless
public class PartsService {

    @PersistenceContext
    EntityManager em;

    public Engine getEngine() {
        Engine e = new Engine();
        em.persist(e);
        return e;
    }

    public SapinVanille getSapin() {
        SapinVanille s = new SapinVanille();
        em.persist(s);
        return(s);
    }
}
```



What happens if there is a shortage of vanilla and a **NullPointerException** is thrown in the constructor?

Transaction scope & JPA

- The default transaction scope for EJB methods is “REQUIRED”. This means that we will have one single transaction for the whole process (and one JPA persistence context).
- **Scenario 1:** no exception thrown. The 3 rows will be inserted when the container commits.

T1	Persistence Context
[CarService] Engine e = getEngine();	PC[t1]= {}
[PartsService] Engine e = new Engine(); em.persist(e);	PC[t1]= {e}
[CarService] SapinVanille s = getSapin();	PC[t1]= {e}
[PartsService] SapinVanille s = new SapinVanille(); em.persist(s);	PC[t1]= {e, s}
[CarService] Car c = new Car(e, s); em.persist(c);	PC[t1]= {e, s, c}

Transaction scope & JPA

- The default transaction scope for EJB methods is “REQUIRED”. This means that we will have one single transaction for the whole process (and one JPA persistence context).
- **Scenario 2:** an exception is thrown in the SapinVanille constructor. No row is added to the database (not even the engine which was successfully persisted).

T1	Persistence Context
[CarService] Engine e = getEngine();	PC[t1]= {}
[PartsService] Engine e = new Engine(); em.persist(e);	PC[t1]= {e}
[CarService] SapinVanille s = getSapin();	PC[t1]= {e}
[PartsService] SapinVanille s = new SapinVanille(); NullPointerException is thrown	PC[t1]= {e}
Transaction is rolled back by the container	

Transaction scope & JPA

- Vanilla shortage should not block the production line!
- The Car constructor is ok with a null value anyway. Let's execute the getSapin() method in its own transaction.

```
@Stateless
public class CarService {

    @PersistenceContext
    EntityManager em;

    @EJB
    PartsService partsService;

    public Car buildCar() {
        Engine e = getEngine();
        try {
            SapinVanille s = getSapin();
        } catch (Exception e) {
            logException(e);
        }
        Car c = new Car(e, s);
        em.persist(c);
    }
}
```

```
@Stateless
public class PartsService {

    @PersistenceContext
    EntityManager em;

    public Engine getEngine() {
        Engine e = new Engine();
        em.persist(e);
        return e;
    }

    @TransactionAttribute(TransactionAttributeType.REQUIRES_NEW)
    public SapinVanille getSapin() {
        SapinVanille s = new SapinVanille();
        em.persist(s);
        return(s);
    }
}
```

Transaction scope & JPA

- **Scenario 1:** no exception thrown. 1 row is committed in the SapinVanille table first, 2 rows are committed in the Engine and Car tables later.

s is not a managed entity

T1	T2	Persistence Context
<code>[CarService]</code> <code>Engine e = getEngine();</code>		PC[t1]= {}
<code>[PartsService]</code> <code>Engine e = new Engine();</code> <code>em.persist(e);</code>		PC[t1]= {e}
<code>[CarService]</code> <code>SapinVanille s = getSapin();</code>		PC[t1]= {e}
	<code>[PartsService]</code> <code>SapinVanille s = new SapinVanille();</code> <code>em.persist(s);</code>	PC[t2]= {s}
	The container commits T2	
<code>[CarService]</code> <code>Car c = new Car(e, s);</code> <code>em.persist(c);</code>		PC[t1]= {e, c}
	The container commits T1	

Transaction scope & JPA

- **Scenario 2:** an exception is thrown in the SapinVanille constructor. No row is committed in the SapinVanille table, BUT 2 rows are committed in the Engine and Car tables!

T1	T2	Persistence Context
[CarService] Engine e = getEngine();		PC[t1]= {}
[PartsService] Engine e = new Engine(); em.persist(e);		PC[t1]= {e}
[CarService] SapinVanille s = getSapin();		PC[t1]= {e}
	[PartsService] SapinVanille s = new SapinVanille(); NullPointerException is thrown	PC[t2]= {}
	The container rolls back T2	
[CarService] Car c = new Car(e, s); em.persist(c);		PC[t1]= {e, c}
	The container commits T1	


Transactions & concurrency control

- If several transactions are processed **concurrently**, unexpected results may occur. There are different strategies and mechanisms for dealing with that.

```
@Stateless
public class TransactionProcessor {

    @EJB
    AccountDAO accountDao;

    public void processTransaction(Transaction t) {
        Account a = accountDao.findById(t.getAccountId());
        long previousBalance = a.getBalance();
        a.setBalance(previousBalance + t.getAmount());
    }
}
```



What happens if another transaction modifies the account balance between these two statements?

Optimistic concurrency control

- In many applications, there is a **high ratio of “read to write” operations** (many transactions read data, few update data). Moreover, there is a “small” likelihood that two concurrent transactions try to update the same data.
- In this case, for performance and scalability reasons, it is often recommended to implement an optimistic concurrency control mechanism.
- The mechanism works as follows:
 - When a program **reads** a record, it gets its “**version number**” (the number of previous updates) in a table column.
 - When it **updates** this record, it makes sure that the version number has not been incremented (this would indicate a conflict with another transaction).
- The developer has to write the logic to execute when a conflict is notified (retry, notify the user, etc.)

Optimistic concurrency control with JPA

- **JPA supports optimistic concurrency control.**
- To use it, the first step is to annotate one field of the entity with the **@version** annotation. JPA will ensure that this value is incremented with every update.
- The second step is to catch the **OptimisticLockException** that may be thrown by JPA when the transaction commits.
- This is where the developer specifies what to do if a conflict has been detected. In some cases, it is possible to immediately and silently retry the transaction.

```
@Entity
public class Account {

    @Id
    long accountId;

    @Version
    long version;
}
```

Pessimistic concurrency control

- When an optimistic concurrency control is not appropriate, then it is possible to implement **pessimistic concurrency control with locks**.
- RDBMS support different types of locks (read lock, write lock).
- When a transaction obtains a **read lock** on a record, it cannot be modified by other transactions. However, it can be read by other transactions.
- When a transaction obtains a **write lock** on a record, it cannot be modified, nor read by other transactions.
- Locking database records **introduce issues**: scalability, performance, deadlocks. It can be tricky to decide when to obtain a lock and for how long.

Pessimistic concurrency control with JPA

- **JPA supports pessimistic concurrency control since version 2.0**
- It is possible to **lock a record** with **em.lock(entity, LOCK_TYPE)**.
- It is also possible to lock a record at the time of retrieval with **em.find(class, id, LOCK_TYPE)**

```
Account a = em.find(Account.class, id);  
em.lock(a, PESSIMISTIC_WRITE);
```

Be aware that we
still have a risk of
stale data here!

```
Account a = em.find(Account.class, id, PESSIMISTIC_WRITE);
```